

Why was nerdctl made?

<https://github.com/containerd/nerdctl>

nerdctl: containerd CTL

- Docker-compatible CLI for containerd (and buildkitd)
`alias docker=nerdctl`
- **Made for facilitating new experiments in the containerd platform**
- The experimental features can be safely turned off for production environments (`export NERDCTL_EXPERIMENTAL=0`)
- Useful for debugging Kubernetes nodes too

Demo

```
$ nerdctl run ...  
$ nerdctl compose up ...
```

Novel features of containerd & nerdctl


- Lazy-pulling with eStargz/SOCI/Nydus/OverlayBD
- Registry-less P2P image distribution with IPFS *
- Image encryption with OIcrypt *
- Image signing with Cosign / Notation *
- “Real” read-only mounts with mount_setattr
- Slirp-less rootless containers with bypass4netns *
- Interactive debugging of Dockerfiles, with buildg *

Lazy-pulling with eStargz/SOCI/Nydus/OverlayBD

- Lazy-pulling: pulling image contents on demand
- No need to pull an entire image
- Several formats are being proposed


Format	Implementation for containerd	Description
eStargz	github.com/containerd/stargz-snapshotter	Optimizes gzip granularity for seek()-ability; Forward compatible with OCI v1 tar.gz
SOCI	github.com/awslabs/soci-snapshotter	Captures a checkpoint of tar.gz decoder state; Forward compatible with OCI v1 tar.gz
Nydus	github.com/containerd/nydus-snapshotter	An alternate image format; Not compatible with OCI v1 tar.gz
OverlayBD	github.com/containerd/overlaybd	Block devices as container images; Not compatible with OCI v1 tar.gz

Lazy-pulling with eStargz/SOCI/Nydus/OverlayBD




Registry-less P2P image distribution with IPFS

```
$ nerdctl pull ipfs://<CID>
```



Registry-based image distribution



IPFS-based image distribution

Image encryption with OCIcrypt

- OCI Image Layers can be encrypted to protect secret files
(OCI Image Config, e.g., env vars, are not encrypted)

```
$ nerdctl image encrypt --recipient=jwe:mypubkey.pem foo example.com/foo:encrypted  
$ nerdctl push example.com/foo:encrypted
```

- Encrypted layers can be decrypted transparently

```
$ nerdctl pull example.com/foo:encrypted
```


Image signing with Cosign / Notation

- Docker CLI supports `DOCKER_CONTENT_TRUST` (Notary v1), but it didn't see wide adoption
- `DOCKER_CONTENT_TRUST` is not implemented for nerdctl
- Instead, nerdctl supports Cosign and Notation (“Notary v2”)

```
$ nerdctl pull --verify=(cosign|notation)
```


“Real” read-only mounts with mount_setattr

- Surprisingly `docker run -v /mnt:/mnt:ro` doesn't really make `/mnt` inside the container read-only
 - Submounts like `/mnt/usbdisk` are writable
- `nerdctl` supports RRO: recursively-read-only mounts
`nerdctl run -v /mnt:/mnt:rro`
 - Implemented using `MOUNT_ATTR_RDONLY` (kernel ≥ 5.12)

Slirp-less rootless containers with bypass4netns

```
$ nerdctl run --label=nerdctl/bypass4netns=1
```

- Bypass slirp4netns (usermode TCP/IP), by using `SECCOMP_IOCTL_NOTIF_ADDFD`
- Even faster than rootful



Interactive debugging of Dockerfiles, with buildg

- Breakpoints for Dockerfiles

```
$ nerdctl builder debug .  
....  
=>  1| FROM ubuntu AS dev  
    2| RUN echo hello > /hello  
    3| RUN echo world > /world  
    4|  
(buildg) break 3  
(buildg) continue  
...  
Breakpoint[0]: reached line: Dockerfile:3  
(buildg) exec /bin/sh  
#
```

Other features


- Namespacing: `nerdctl --namespace=<NS> ps`
- Connecting a container to multiple networks at once:
`nerdctl run --net foo --net bar ...`
- AppArmor for rootless containers
(Although still needs `sudo nerdctl apparmor load`)
- FreeBSD containers
- Windows containers
- And more...

FAQ: Why not ctr?

- ctr: a CLI included in containerd
- ctr is often misbelieved to be the standard CLI for containerd
- But ctr is just a debugging tool written for containerd developers
- ctr lacks lots of high-level features that are present in Docker and nerdctl
 - Exposing ports (`nerdctl run -p <PORT>`)
 - Re-startable containers w/ networking (`nerdctl run --restart=always`)
 - Support for `~/.docker/config.json`, `docker-credential-ecr-login`, etc.
 - Reading logs of background containers (`nerdctl logs`)

FAQ: Why not crictl?

- crictl is similar to ctr but uses a different API
- crictl has similar restrictions as ctr



- CRI API has very tight scope compared to the containerd native API
- Negotiation for modifying the CRI API is relatively hard

FAQ: Why not Docker?

- Docker doesn't/didn't use the containerd API for image management
 - Can't/Couldn't be used for experimenting lazy-pulling, etc.
 - Work is in progress toward Docker v24+
- Release cycle is/was slow
 - Docker v19.03: Jul 2019
 - Docker v20.10: Dec 2020
 - Docker v23: Feb 2023
- Not trying to “defeat” Docker; Experiments in nerdctl are being ported to Docker too

FAQ: Why not Podman/CRI-O ?

- Podman doesn't work as a Kubernetes runtime
- CRI-O doesn't support non-CRI API, and lacks Docker-like CLI
 - `podman ps`, etc. works for CRI-O too, but still Podman != CRI-O
- Not trying to “defeat” Podman/CRI-O either;
Experiments in containerd/nerdctl are being ported to Podman/CRI-O too
 - e.g., eStargz

Getting started

- <https://github.com/containerd/nerdctl/releases>
- `nerdctl-full-<VERSION>-linux-amd64.tar.gz` contains all the dependencies (containerd, runc, ...)

```
$ sudo systemctl enable --now containerd  
$ sudo nerdctl run -d --name nginx -p 80:80 nginx:alpine
```

```
$ containerd-rootless-setup.sh install  
$ nerdctl run -d --name nginx -p 8080:80 nginx:alpine
```

Getting started

Lima

- macOS users should try [Lima](#)
 - Lima: Linux Machine, originally designed as “nerdctl Machine”
 - CNCF Sandbox Project
 - Automatic port forwarding & host filesystem mounting

```
$ brew install lima  
$ alias nerdctl=nerdctl.lima  
$ nerdctl run -d --name nginx -p 80:80 nginx:alpine
```

- nerdctl is also included in [AWS Finch](#), [Colima](#), [Rancher Desktop](#), ...



Wrap-up

- nerdctl = containerd CTL
- Made for facilitating new experiments in the containerd platform

- Lazy-pulling with eStargz/SOCI/Nydus/OverlayBD
- Registry-less P2P image distribution with IPFS
- Image encryption with OClcrypt
- Image signing with Cosign / Notation
- “Real” read-only mounts with mount_setattr
- Slirp-less rootless containers with bypass4netns
- Interactive debugging of Dockerfiles, with buildg
- And more!