

# Kubernetes in Rootless Podman

Akihiro Suda, NTT


# Rootless containers

- A technique to run container runtimes as a non-root user
- Available for LXC, Docker, Podman, containerd, etc.
- Mitigates potential vulnerabilities of container runtimes
  - Even if it gets compromised, it will not affect files and processes owned by other user IDs
  - Less chance of having stealth malware, as the kernel, firmware, etc., are protected
  - No ARP spoofing/DNS spoofing in the physical network

<https://blog.aquasec.com/dns-spoofing-kubernetes-clusters>

# Rootless containers

- Implemented by using User Namespaces
  - A feature of the Linux kernel
  - Maps the root in the UserNS to a non-root user outside the UserNS
  - `dnf`, `apt-get`, etc. just work, because they think they are running as the root



# Rootless Kubernetes

- Began in 2018 <https://twitter.com/AkihiroSuda/status/1019570064385642498>
  - As old as Rootless Docker (pre-release at that time) and Rootless Podman
- The changes to Kubernetes was merged in Kubernetes v1.22 (Aug 2021)
  - Feature gate: "KubeletInUserNamespace" (Alpha)

# KubeletInUserNamespace feature gate

- Slightly misnomer; it refers to running all the node components (kubelet, kube-proxy, CRI, CNI, OCI) in UserNS
- Root-in-UserNS is similar to the root, but has no permission for:
  - some sysctls
  - dmesg
- The feature gate allows ignoring these permission errors

<https://github.com/search?q=repo%3Akubernetes%2Fkubernetes%20KubeletInUserNamespace&type=code>

# How to run Rootless Kubernetes

The easiest way to run Rootless Kubernetes today is to wrap a Kubernetes node in a Rootless container (such as Rootless Podman)

- kind
- minikube
- Usernetes (Gen2)

# kind (Kubernetes in Docker)

- <https://kind.sigs.k8s.io/>
- The most typical way to run Kubernetes in Docker (and in Podman)
- Supports multi-node, but only on a single host
  - 1 kind container = 1 Kubernetes node
- Not intended to be used for production environments

# kind (Kubernetes in Docker): Usage

- A few of steps needs to be executed by the root
  - These steps are needed for minikube, Usernetes, etc. too

Needs cgroup v2  
(RHEL >= 9, etc.)

```
# Allow limiting CPU, memory, etc. via cgroups
cat <<EOF | sudo tee \
/etc/systemd/system/user@.service.d/delegate.conf
[Service]
Delegate=cpu cpuset io memory pids
EOF
sudo systemctl daemon-reload
```



# kind (Kubernetes in Docker): Usage

- A few of steps needs to be executed by the root
  - These steps are needed for minikube, Usernetes, etc. too

```
# Load extra kernel modules
cat <<EOF | sudo tee /etc/modules-load.d/iptables.conf
ip6_tables
ip6table_nat
ip_tables
iptable_nat
EOF
systemctl restart systemd-modules-load.service
```

# kind (Kubernetes in Docker): Usage

- <https://kind.sigs.k8s.io/docs/user/rootless/>

```
export KIND_EXPERIMENTAL_PROVIDER=podman
```

```
kind create cluster
```

```
kubectl get pods -A
```

# minikube

- <https://minikube.sigs.k8s.io/>
- Originally designed for running Kubernetes in VM
- Supports kind-like mode too

# minikube: Usage

- <https://minikube.sigs.k8s.io/docs/drivers/podman/>

```
minikube config set rootless true
```

```
minikube start --driver=podman --container-runtime=crio
```

```
kubectl get pods -A
```

- Make sure to set "rootless" property, otherwise minikube executes podman with sudo

- <https://github.com/rootless-containers/usernetes>
- Rootless Kubernetes, since 2018
  - **Gen 1** (2018-2023): “The hard way”
  - **Gen 2** (2023-): depends on Rootless (Docker|Podman|nerdctl) for simplicity
- Supports real multi-node clusters with VXLAN

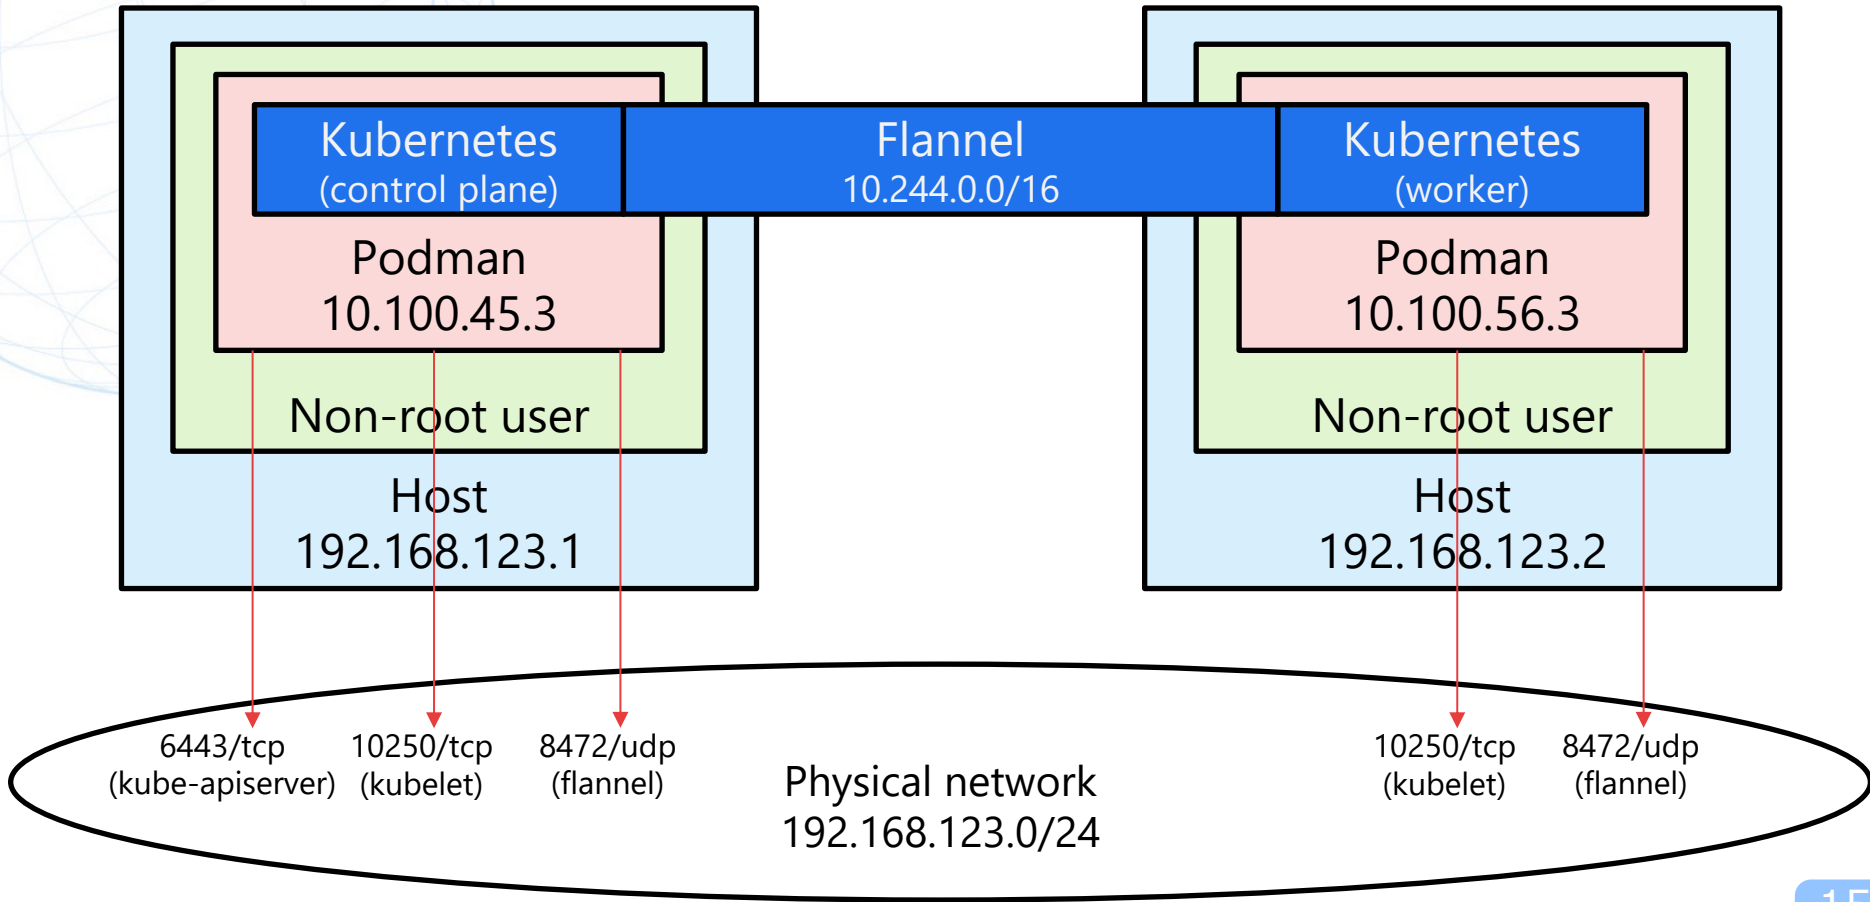
# Usernetes: Gen 1 vs Gen 2

"The hard way"

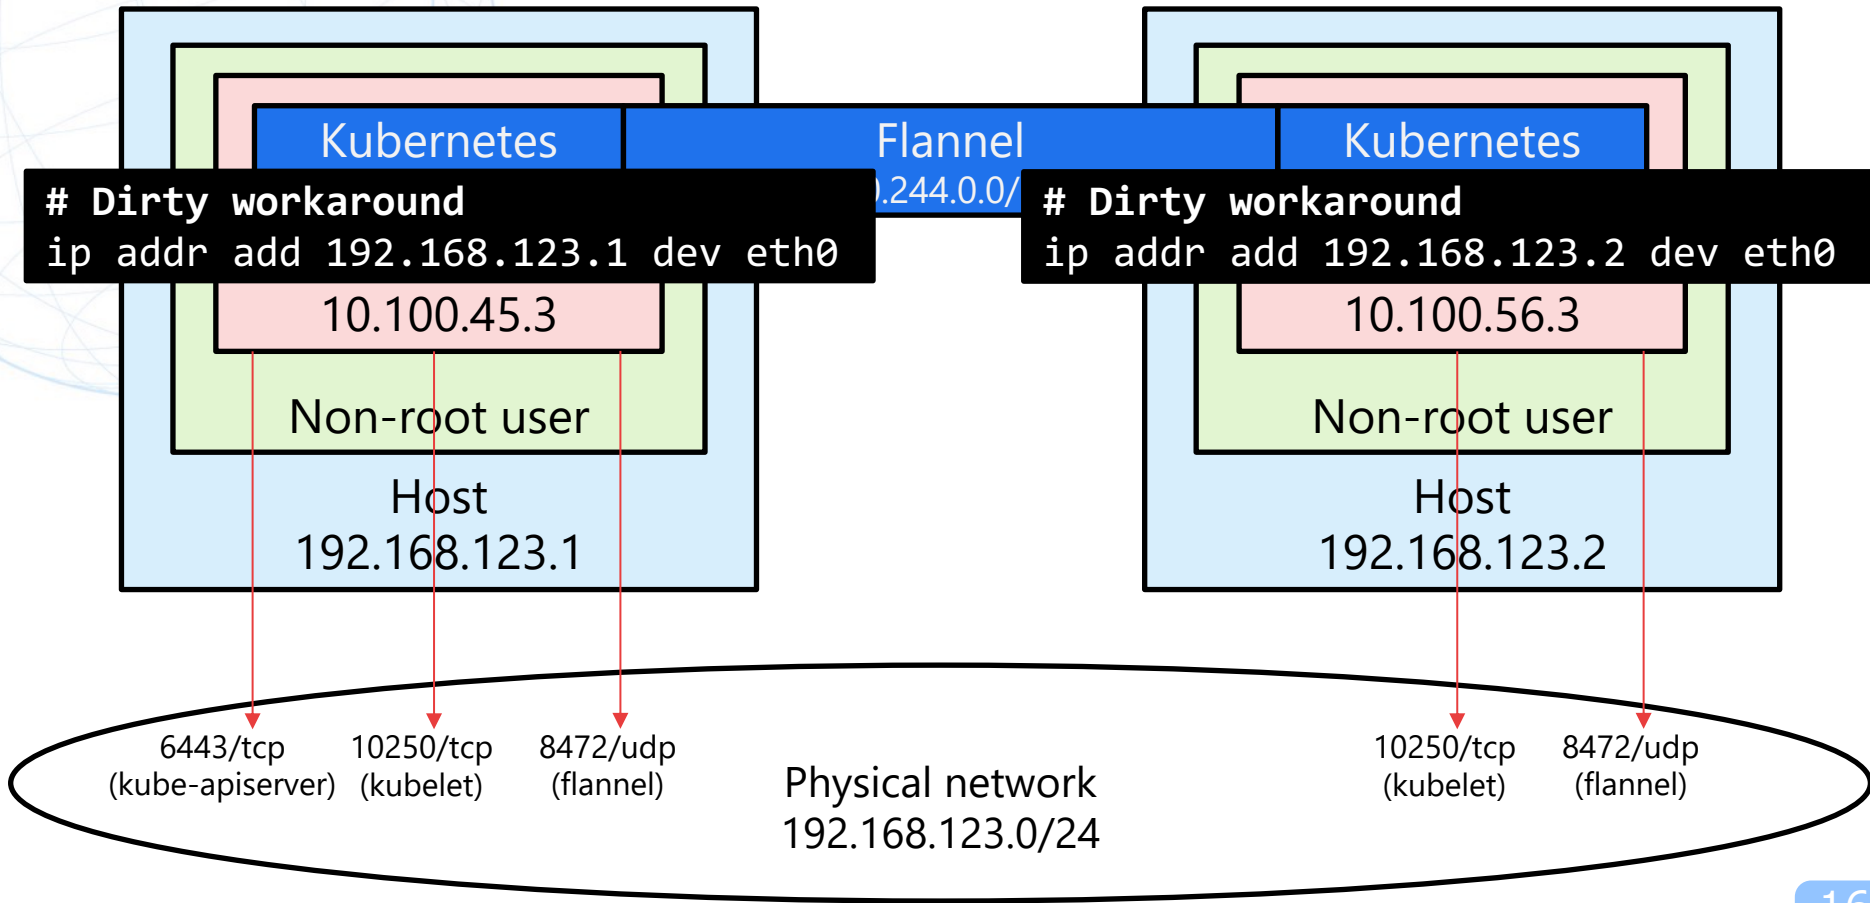
Similar to `kind` and minikube,  
but supports real multi-node

	Gen 1 (2018-2023)	Gen 2 (2023-)
<b>Host dependency</b>	RootlessKit	Rootless Docker, Rootless Podman, or Rootless nerdctl (contaiNERD CTL)
<b>Supports kubeadm</b>	No	Yes
<b>Supports multi-node (multi-host)</b>	Yes, but practically No, due to complexity	Yes
<b>Supports hostPath volumes</b>	Yes	Yes, for most paths, but needs an extra config

# Usernetes (Gen 2): How it works



# Usernetes (Gen 2): How it works





# Usernetes (Gen 2): Usage

Set `CONTAINER\_ENGINE=podman`  
if multiple container engines are  
installed on the host

## # Bootstrap the first node

```
make up  
make kubeadm-init  
make install-flannel
```

## # Enable kubectl

```
make kubeconfig  
export KUBECONFIG=$(pwd)/kubeconfig  
kubectl get pods -A
```


## # Multi-node

```
make join-command  
scp join-command another-host:~/usernetes  
ssh another-host make -C ~/usernetes up kubeadm-join
```

# Future works

Multi-tenancy using multiple user IDs and multiple TCP ports

- A single host will be able to join multiple clusters



# Future works

Promote “KubeletInUserNamespace” gate from alpha to beta (and then GA)

- The blocker was how to test the gate in the upstream CI
- WIP: <https://github.com/kubernetes/test-infra/pull/31085>
  - Spawns rootless `kind` machines using Google Compute Engine

# Future works

Eliminate the overhead of user-mode TCP/IP  
(slirp4netns, RootlessKit, and pasta)

- POC: <https://github.com/rootless-containers/bypass4netns>
- Captures socket-related syscalls in containers using seccomp\_unotify(2), and replaces the socket FDs with ones that are created in the host network namespace
- Unsolved question: how to support VXLAN?  
VXLAN is implemented in the kernel, so VXLAN calls cannot be captured with seccomp\_unotify(2)

# Future works

Support running okd (OpenShift) in Rootless Podman

- Help wanted from the OpenShift community