

# Usernetes Gen2

## Kubernetes in Rootless Docker, with Multiple Nodes

<https://github.com/rootless-containers/usernetes>

# Usernetes Gen2

Kubernetes in Rootless Docker, with Multiple Nodes  
Podman, and nerdctl

<https://github.com/rootless-containers/usernetes>

# [Introduction] Rootless Containers


- Puts container runtimes (as well as containers) in a user namespace
  - UserNS: Linux kernel's feature that maps a non-root user to a fake root (the root privilege is limited inside the namespace)
- Can mitigate potential vulnerabilities of the runtimes
  - No access to read/write other users' files
  - No access to modify the kernel (e.g., to inject invisible malware)
  - No access to modify the firmware
  - No ARP spoofing
  - No DNS spoofing
- Also useful for shared hosts (High-performance Computing, etc.)
  - Works with GPU too

e.g., runc breakout  
CVE-2024-21626  
(2024-01-31)

# [Introduction] User namespaces

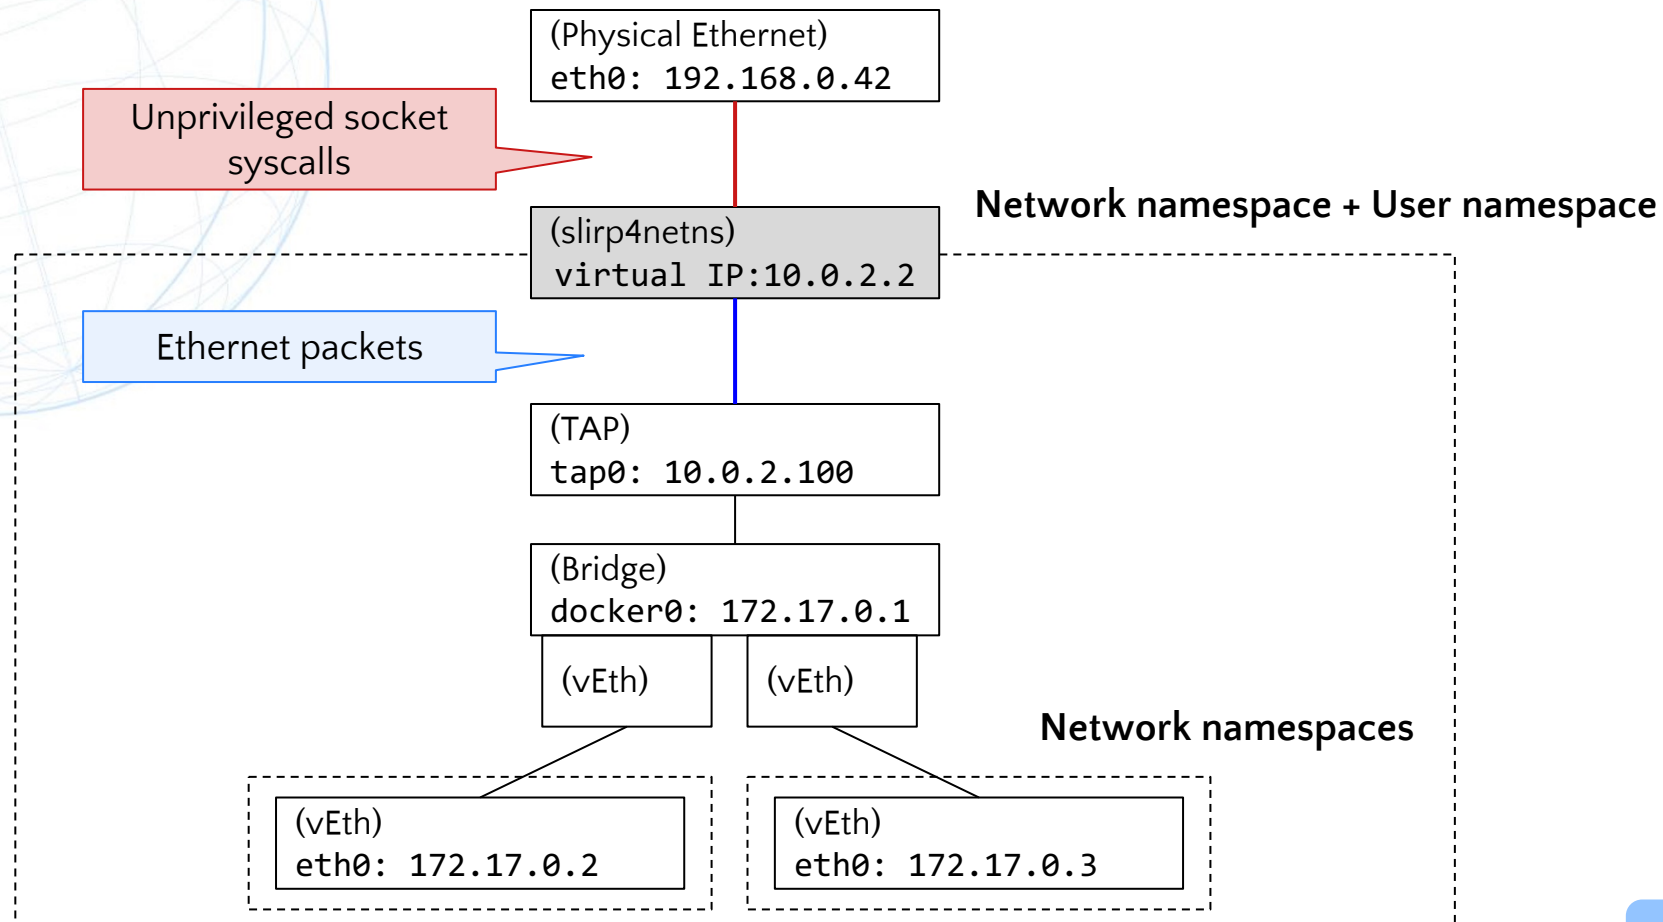
- Linux kernel's feature to remap UIDs and GIDs

```
# /etc/subuid  
1000:100000:65536
```



- UID=1000 gains **fake** root privileges (UID=0) that are enough to create containers
- The privileges are limited inside the namespace
- No privilege for setting up vEth pairs with “real” IP addresses; user mode TCP/IP (e.g., slirp4netns) is used instead
- Also notorious as the culprit of the several kernel CVEs, but at least it is more secure than just running everything as the root
  - Ubuntu 24.04 disables UserNS by default with the allowlist (AppArmor profiles)

# [Introduction] Network namespaces



- **Usernetes:** Rootless Kubernetes, since 2018  
<https://github.com/rootless-containers/usernetes>
- As old as Rootless Docker (pre-release at that time) and Rootless Podman
- The changes to upstream was merged in Kubernetes v1.22 (2021)
  - Feature gate: `KubeletInUserNamespace` (Alpha)
  - The feature gate is also used by `kind`, `minikube`, `k3s`, etc.
- The first generation (“Gen1”, 2018–2023) of Usernetes didn’t gain much popularity due to its complexity (“The Hard Way”)

# KubeletInUserNamespace feature gate

- The gate is slightly misnomer; as it requires CRI, OCI, CNI, and kube-proxy to be in the same UserNS too
- Quite “boring” gate to allow trivial permission errors  
<https://github.com/search?q=repo%3Akubernetes%2Fkubernetes%20KubeletInUserNamespace&type=code>
  - `dmesg`
  - `sysctl -w vm.overcommit_memory`
  - etc.
- The UserNS has to be created by an external runtime
  - **Usernetes Gen1:** RootlessKit
  - **Usernetes Gen2:** Rootless Docker/Podman/nerdctl
  - LXD/Incus can be used too

# Usernetes Gen 1 vs Gen 2

"The Hard Way"

Similar to `kind` and minikube,  
but supports real multi-node

	Gen 1 (2018-2023)	Gen 2 (2023-)
<b>Host dependency</b>	RootlessKit	Rootless Docker, Rootless Podman, or Rootless nerdctl (contaiNERD CTL)
<b>Supports kubeadm</b>	No	Yes
<b>Supports multi-node</b>	Yes, but practically No, due to complexity	Yes
<b>Supports hostPath volumes</b>	Yes	Yes, for most paths, but needs an extra config



# File layout

Everything is just a plain text file,  
for ease of customization

- **Makefile**
  - Defines targets like `make up` to wrap `docker compose up`, etc.
- **Dockerfile**
  - `FROM kindest/node` (kind's node image) with a few additional `ADD` and `RUN`
- **docker-compose.yaml**
  - Just defines a single node container
  - Currently, node ports, etc. have to be statically defined here
- **kubeadm-config.yaml**
  - Configures feature gates, CIDRs, TLS SANs, etc.

## # Bootstrap the first node

```
make up  
make kubeadm-init  
make install-flannel
```

## # Enable kubectl

```
make kubeconfig  
export KUBECONFIG=$(pwd)/kubeconfig  
kubectl get pods -A
```

## # Multi-node

```
make join-command  
scp join-command another-host:~/usernetes  
ssh another-host make -C ~/usernetes up kubeadm-join  
make sync-external-ip
```

# Multi-node Network

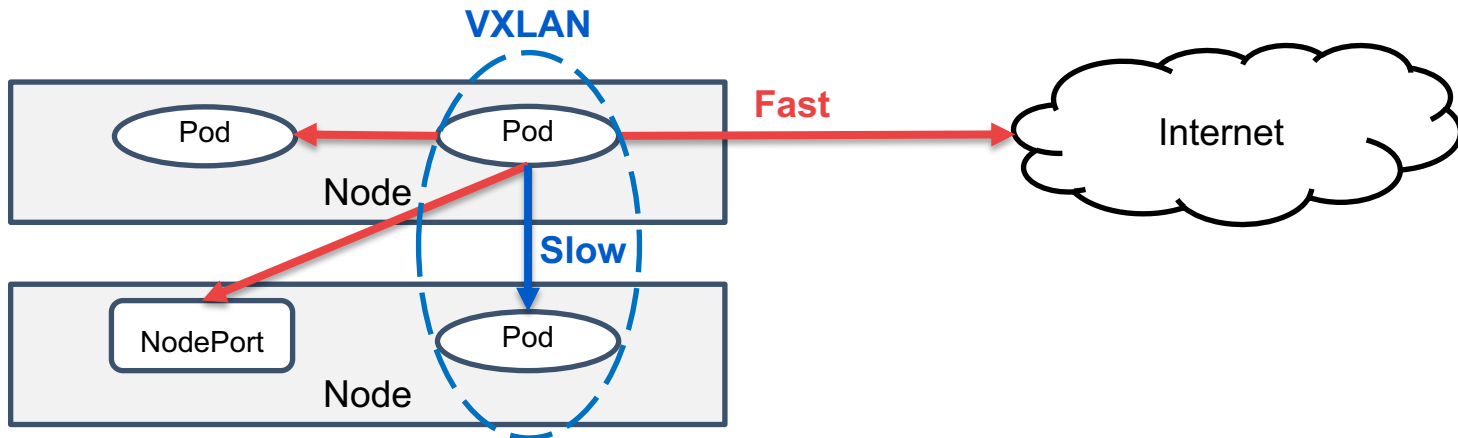
- VXLAN is known to work
  - Just `kubect1 -f kube-flannel.yaml`
- “External IP” is used, as the containerized kubelet’s IP is not accessible from other nodes
  - `kubelet` is launched with `--cloud-provider=external`
  - `node.status.addresses` is dynamically patched with `kubect1 patch node`
  - `node` is also annotated with `flannel.alpha.coreos.com/public-ip-overwrite`
  - UDP checksums are recomputed with `ethtool -K flannel.1 tx-checksum-ip-generic off`

# Experimental: network acceleration

- Bypass4netns allows bypassing slirp4netns to eliminate the overhead caused by the usermode TCP/IP  
<https://github.com/rootless-containers/bypass4netns>
- Captures socket syscalls inside the NetNS, reconstructs the FDs outside the NetNS, and replaces the FDs inside the NetNS, using `seccomp_unotify(2)`
- As fast as the host network (e.g., 1.28 Gbps vs 49.9 Gbps)
- Bypass4netns supports both `connect(2)` and `bind(2)`, but Usernetes only supports accelerating `connect(2)` currently
  - `bind(2)` is already fast anyway
- Available for nerdctl

# Experimental: network acceleration

- Pod-to-Pod communications across multiple nodes are not accelerated yet
  - VXLAN packets are generated by the kernel itself and cannot be intercepted via `seccomp_unotify(2)`
  - NodePorts can be still accelerated, as it does not incur VXLAN packets



# Experimental: network acceleration

- iperf3 (TCP) benchmark across multiple nodes

	slirp4netns	bypass4netns
<b>Pod → Pod (same node)</b>	37.6 Gbps	37.6 Gbps
<b>Pod → Pod (different node)</b>	1.40 Gbps	1.41 Gbps
<b>Pod → NodePort (same node)</b>	1.28 Gbps	<b>49.9 Gbps</b>
<b>Pod → NodePort (different node)</b>	1.47 Gbps	<b>9.53 Gbps</b>
<b>Host → NodePort (same node)</b>	50.2 Gbps	49.4 Gbps
<b>Host → NodePort (different node)</b>	9.53 Gbps	9.52 Gbps

IaaS: Amazon EC2 (m7i.2xlarge)

Versions: Ubuntu 22.04, nerdctl v2.0.0-beta.4, bypass4netns v0.4.1, Usernetes Gen2-v20240410.0 (Kubernetes v1.29)

<https://github.com/rootless-containers/usernetes/discussions/329>

# Future works

- Integrate bypass4netns into Docker and Podman too
- Support accelerating Pod-to-Pod communications across different nodes, perhaps with a sidecar proxy that would forward packets to NodePorts
- Support dynamic port forwarding
  - Ports are currently statically defined in `docker-compose.yaml`
  - If `docker container update` could support modifying port forwards, Usernetes could just watch Kubernetes service events and update the Docker ports accordingly  
<https://github.com/docker/cli/issues/5013>
- Help other Kubernetes distributions to support rootless
  - k3s has been supporting rootless since 2019, but still lacks support for multi-node setup
  - Are Podman folks interested in running OKD inside Rootless Podman?