


containerd v2.0



An industry-standard container runtime


PromCon

North America 2021

- CNCF Graduated since 2019
- Adopted by well-known enterprise Kubernetes services and distributions
- Designed to be pluggable to support cutting-edge features (remote snapshotters, WASM runtimes, ...)


Runtime plugins

- **runc**: the regular runtime for Linux
- **runhcs**: Windows
- **runj**: FreeBSD jail
- **runwasi**: WASM
- **kata**: VM
- **runsc (gvisor)**: ptrace sandbox, etc.



PromCon
North America 2021

Snapshotter plugins



Regular snapshotters: **overlayfs**, **btrfs**, **zfs**, **devmapper**, ...

DockerCon
North America 2021

“Remote” snapshotters support pulling image contents on demand to shorten the container startup time

- **stargz**: Forward compatible with OCI v1 tar.gz images
- **nydus**: Uses an alternate image format
- **overlaybd**: Uses block devices as container images

containerd v2.0

- The focus is on cleaning up deprecated features (runtime v1 API, AUFS snapshotter, non-sandboxed pods...)
- Huge improvements on the support for User Namespaces
 - UserNS maps the root user in containers to a non-root user on the host to mitigate potential container break-out
 - Overhead for chowning is now eliminated




Snapshotter is overlayfs.

Image	Size	Inodes	Chown without metacopy	Chown with metacopy	IDMapped Mount
tensorflow/tensorflow:2.13.0	1.6GiB	37460	2m58.18635105s	1.425115168s	21.969618ms
library/node:20.11.1	1.1GiB	33210	2m40.427234391s	1.047708275s	17.713161ms
library/ubuntu:22.04	84MiB	3517	14.986225994s	101.244062ms	9.268065ms

Deprecation Warning


- ctr deprecations list
 - Supported in v1.6.27+, v1.7.12+
 - For each warning, you can see
 - What is being removed
 - A suggestion for migration
 - When it was last used



PromCon
North America 2021

[Deprecation Warnings in containerd - Getting Ready for 2.0!](#)

nerdctl: contaiNERD CTL




- <https://github.com/containerd/nerdctl>
- Same UI/UX as the docker CLI (including Compose)
- Made for facilitating new experiments in the containerd platform (e.g., stargz, fast rootless with bypass4netns)
- Useful for debugging Kubernetes nodes too

```
$ nerdctl run hello-world  
$ nerdctl compose up
```

PromCon
North America 2021

nerdctl v2.0




PromCon

North America 2021

- Rootless networking is refactored to support “detach-netns”
 - Faster pull/push
 - Previously limited to less than 10Gbps due to user-mode TCP/IP;
Now as fast as rootful
 - Containers can be accelerated too with bypass4netns (experimental)
 - Proper support for `--net=host`
 - Proper support for localhost registries

nerdctl v2.0




PromCon
North America 2021

- Support running systemd in a container without `--privileged`
 - "Plain old VM"-like user experience
 - Often considered to be an anti-pattern
 - Useful for testing, etc.

```
$ nerdctl run --systemd=true ...
```

- Massive refactoring and testing


Join the community!



PromCon
North America 2021

- **GitHub:** <https://github.com/containerd>
- **Slack:** <https://slack.cncf.io/>
(Channels: `#containerd` `#containerd-dev`)
- **X (former Twitter):** [@containerd](https://twitter.com/containerd)
- **Bluesky:** [@containerd.dev](https://bsky.app/profile/containerd.dev)

Visit us again on Friday (2:55pm-3:30pm)



What's Going on in the Containerd Neighborhood?

PyCon

North America 2021

Phil Estes, AWS

Samuel Karp, Google

Akihiro Suda, NTT

Michael Brown, IBM

Kirtana Ashok, Microsoft



Hyatt Regency | Level 4 | Regency Ballroom A


Learn about containerd 2.0

What containerd 2.0 means for you

Samuel Karp, Google


Friday, 11:55am - 12:30pm

Salt Palace | Level 2 | 254 B



PromCon
North America 2021





PromCon
North America 2021

Advanced topics

New features, newly stable features, defaults in v2.0



PromCon
North America 2021


- Transfer service **NOW STABLE**
- Sandbox service (and sandboxed CRI) **NOW STABLE**
- Faster image extraction with igzip **NEW**
- Improved OTEL configuration **NEW**
- NRI enabled by default **NEW**
- Image verifier plugins **NEW**
- Plugin introspection **NEW**
- CDI enabled by default **NEW**
- CRI support for user namespaces **NEW**

Highlight: Node Resource Interface (NRI)



- Akin to a mutating webhook, but for container configuration
 - Middleware between CRI and OCI
- Use cases
 - Injection (devices, network devices, OCI hooks)
 - Resource modification/management (ulimits, topology/NUMA, advanced QoS, SGX memory)
 - Policy enforcement
- Plugins can run in containers or as system services
- Enabled by default
- Community plugins
 - <https://github.com/containerd/nri/tree/main/plugins>
 - <https://github.com/containerd/nri-plugins>

Highlight: Image verifier plugins



PromCon

North America 2021

- Exec-based plugins containerd invokes during image pull
- Policy enforcement use-cases
 - Container image signature verification
 - Trust for particular signers
 - Allow only specific registries/repositories
- Integrated with the Transfer service (not supported for legacy pulls)